



Databases & XML (11) 02.05.2016



Time	Agenda
8.30	Polyglot persistence
8.45	Quiz
9.00	Choosing your database
9.20	Mongo vs SQL
9.35	API exercise
10.00	Pause
10.30	Schemaless
10.35	CAP and BASE
10.50	Groups for interdisciplinary project
11.00	Mongo exercises
12.00	Lunch
12.30	About the exam
12.45	Test exam
13.45	Experiences with the exam

## Agenda

---

- [Internet Week Denmark Festival](#) – May 9th – 13th in Aarhus
- [Open Data](#) – brug offentlige data til innovation – May 12th

# Internet Week Denmark

---

BUSINESS ACADEMY  
AARHUS






Figure 13.1 *Use of RDBMS for every aspect of storage for the application*

## Polyglot Persistence (RDBMS)

---

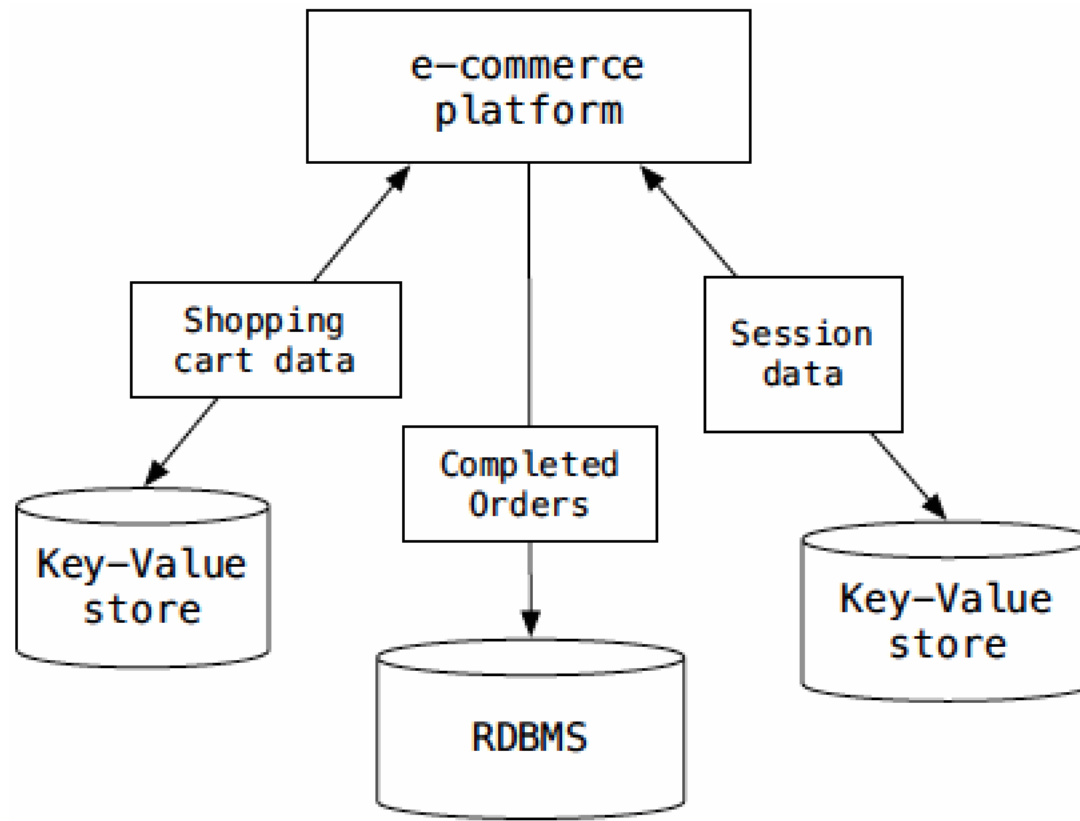


Figure 13.2 *Use of key-value stores to offload session and shopping cart data storage*

## Polyglot Persistence (Polyglot)

---

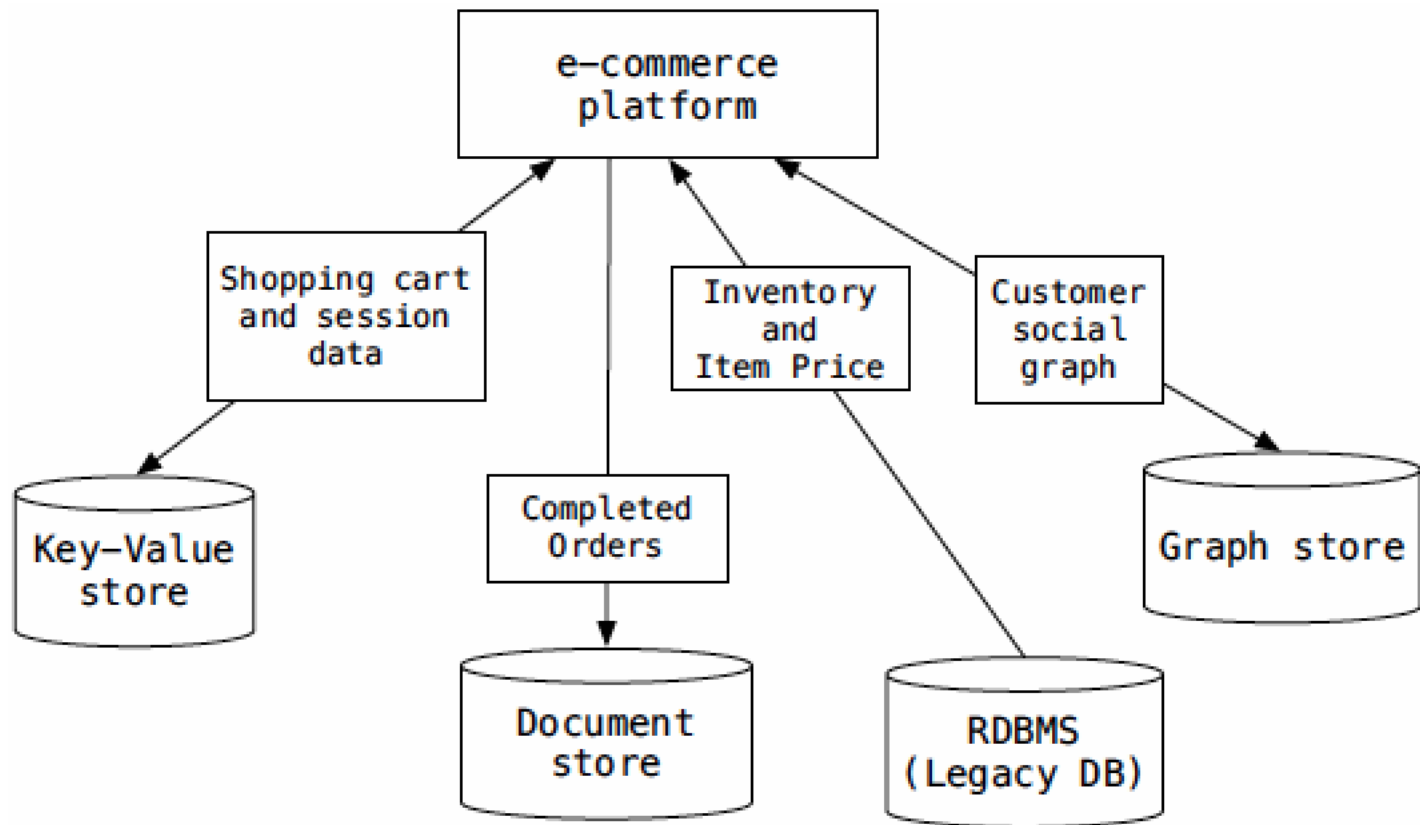


Figure 13.3 *Example implementation of polyglot persistence*

## Polyglot Persistence (Polyglot)

---

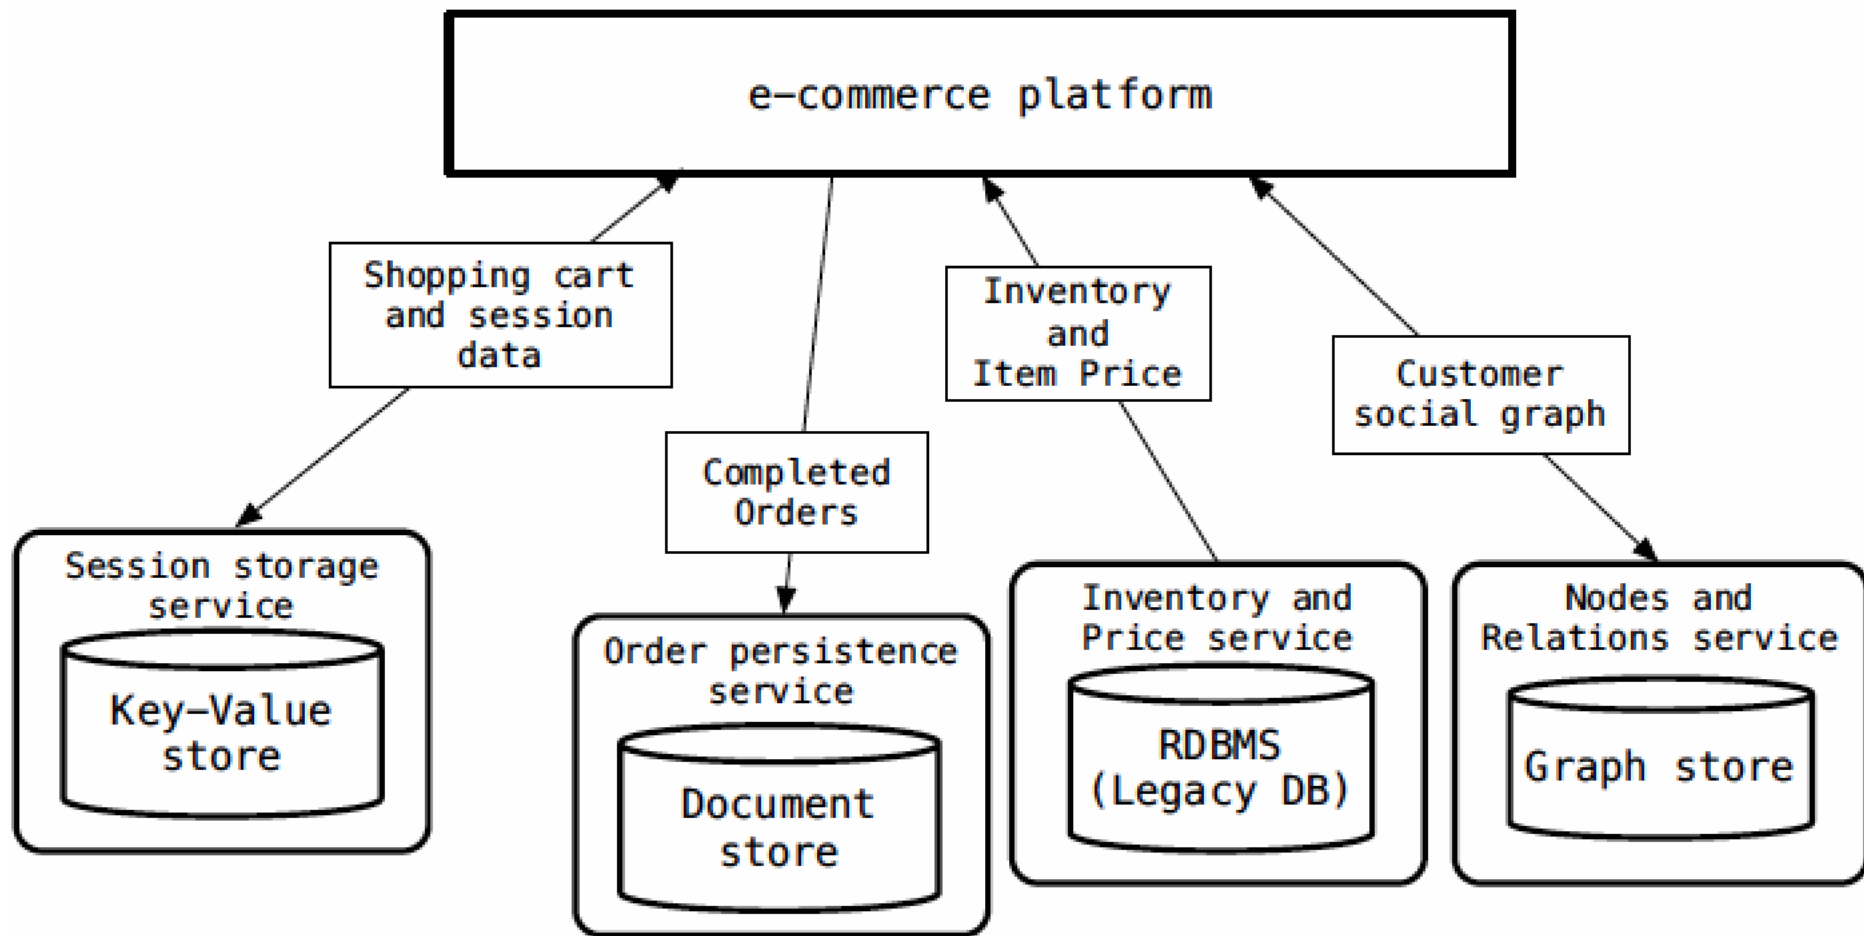


Figure 13.5 *Using services instead of talking to databases*

## Polyglot Persistence (Polyglot)

---

## 13.8 Key Points

- Polyglot persistence is about using different data storage technologies to handle varying data storage needs.
- Polyglot persistence can apply across an enterprise or within a single application.
- Encapsulating data access into services reduces the impact of data storage choices on other parts of a system.
- Adding more data storage technologies increases complexity in programming and operations, so the advantages of a good data storage fit need to be weighed against this complexity.

### Polyglot Persistence (Key Points)

---





## 13.8 Key Points

- Polyglot persistence is about using different data storage technologies to handle varying data storage needs.
- Polyglot persistence can apply across an enterprise or within a single application.
- Encapsulating data access into services reduces the impact of data storage choices on other parts of a system.
- Adding more data storage technologies increases complexity in programming and operations, so the advantages of a good data storage fit need to be weighed against this complexity.

### Polyglot Persistence (Key Points)

---



## 13.8 Key Points

- Polyglot persistence is about using different data storage technologies to handle varying data storage needs.
- Polyglot persistence can apply across an enterprise or within a single application.
- Encapsulating data access into services reduces the impact of data storage choices on other parts of a system.
- Adding more data storage technologies increases complexity in programming and operations, so the advantages of a good data storage fit need to be weighed against this complexity.

### Polyglot Persistence (Key Points)

---





## 13.8 Key Points

- Polyglot persistence is about using different data storage technologies to handle varying data storage needs.
- Polyglot persistence can apply across an enterprise or within a single application.
- Encapsulating data access into services reduces the impact of data storage choices on other parts of a system.
- Adding more data storage technologies increases complexity in programming and operations, so the advantages of a good data storage fit need to be weighed against this complexity.

### Polyglot Persistence (Key Points)

---





<http://martinfowler.com/bliki/PolyglotPersistence.html>

<http://memeagora.blogspot.dk/2006/12/polyglot-programming.html>

<http://www.techrepublic.com/article/developers-are-calling-it-quits-on-polyglot-programming/>

## Polyglot and Full-Stack

---





Quiz time!



- Programmer Productivity
- Data-Access Performance
- Sticking with the Default

## Choosing your database

---



- Look at what your software will need to do.
- Run through the current features and see if and how the data usage fits.
- As you do this, you may begin to see that a particular data model seems like a good fit.
- That closeness of fit suggests that using that model will lead to easier programming.

## Programmer Productivity

---



Polyglot persistence is about using multiple data storing solutions.

It may be that you'll see different data storage models fit different aspects of your data. Using multiple databases is inherently more complex than using a single store, but the advantages of a good fit in each case may be better overall.

## Polyglot Persistence

---





- An aggregate-oriented database may be very fast for reading or retrieving aggregates compared to a relational database where data is spread over many tables.
- Test your potential databases' performance in the scenarios that matter to you. Reasoning about how a database may perform can help you build a short list, but the only way you can assess performance properly is to build something, run it, and measure it.
- You can't build your actual system, so you need to build a representative subset. Choose scenarios that are the most common, the most performance-dependent, and those that don't seem to fit your database model well. The latter may alert you to any risks outside your main use cases.

## Data-Access Performance

---



- There are many cases where you're better off sticking with the default option of a relational database:
  - They are well known; you can easily find people with experience of using them
  - They are mature, so you are less likely to run into the rough edges of new technology
  - There are lots of tools that are built on relational technology that you can take advantage of
  - You don't have to deal with the political issues of making an unusual choice – picking a new technology will always introduce a risk of problems should things run into difficulties
- To choose a NoSQL database you need to show a real advantage over relational databases for your situation.
- There are many cases where it is advantageous to use NoSQL – but not “all” or even “most”.

## Sticking with the Default

---



Break – 5 minutes

---

# RDBMS vs NoSQL

---

## RDBMS

- Structured and organized data
- Structured query language (SQL)
- Data and its relationships are stored in separate tables.
- Data Manipulation Language, Data Definition Language
- Tight Consistency
- BASE Transaction

## NoSQL

- Stands for Not Only SQL
- No declarative query language
- No predefined schema
- Key-Value pair storage, Column Store, Document Store, Graph databases
- Eventual consistency rather ACID property
- Unstructured and unpredictable data
- CAP Theorem
- Prioritizes high performance, high availability and scalability

## NoSQL vs RDBMS

---



## **Vertical scaling**

To scale vertically (or scale up) means to add resources within the same logical unit to increase capacity. For example to add CPUs to an existing server, increase memory in the system or expanding storage by adding hard drive.

## **Horizontal scaling**

To scale horizontally (or scale out) means to add more nodes to a system, such as adding a new computer to a distributed software application. In NoSQL system, data store can be much faster as it takes advantage of “scaling out” which means to add more nodes to a system and distribute the load over those nodes.

# Scalability

---



SQL Terms/Concepts	MongoDB Terms/Concepts
database	<a href="#">database</a>
table	<a href="#">collection</a>
row	<i>document</i> or <i>BSON</i> document
column	<a href="#">field</a>
index	<a href="#">index</a>
table joins	embedded documents and linking
primary key	<a href="#">primary key</a>
Specify any unique column or column combination as primary key.	<a href="#">In MongoDB, the primary key is automatically set to the <code>_id</code> field.</a>
aggregation (e.g. group by)	aggregation pipeline

## SQL vs NoSQL

---



“What answers do I have?”





“What questions do I have?”

Data modelling

---






<http://blog.michaelckennedy.net/2010/04/29/mongodb-vs-sql-server-2008-performance-showdown/>

## MongoDB vs. SQL Server





Scalability/Functionality

---

# APIs

So far we have looked a little bit into the APIs of:

- Youtube
- OMDb (movie database)

Look online and find an API that you are interested in.

Output some interesting information from that API on some webpage (e.g. your videowatcher page).

## Exercise

---



Break 30 minutes

---

# Spot the data types

```
1  {
2    "product": "t-shirt",
3    "price": 12.50,
4    "popular": true,
5    "stock": 300,
6    "tags": ["organic", "printed"],
7    "reviews": null,
8    "variants":
9      [
10       {"variant": "Black Night", "image": "black.jpg"},
11       {"variant": "Blue Morning", "image": "blue.jpg"}
12     ]
13  }
```

## Schemaless


---



CAP og BASE

---

Availability



Consistency

Partition Tolerance

CAP

---



# Consistency

All clients see the same data after execution of operation

# Availability

Always on – no downtime

# Partition Tolerance

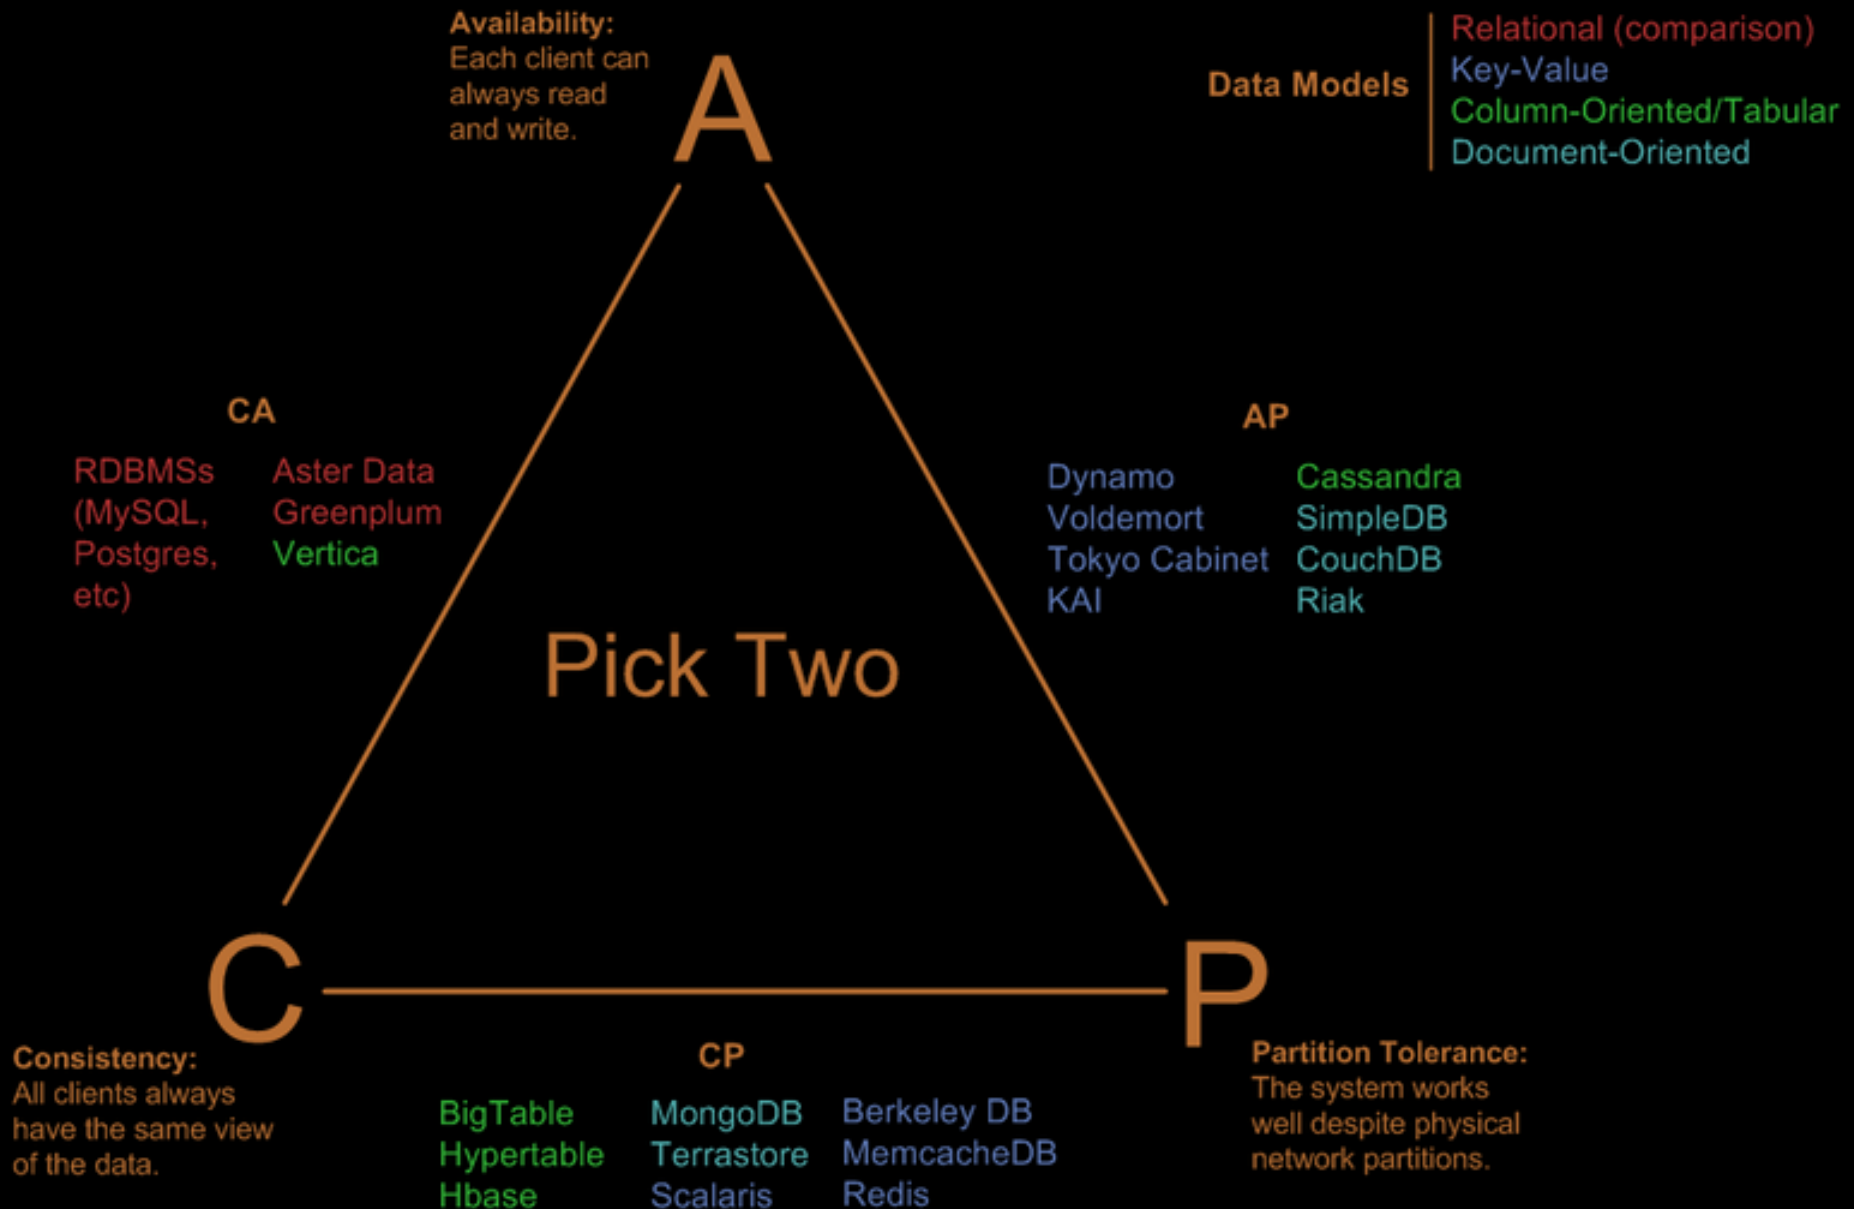
Continues to function in case of failure in a node

CAP Theorem

---



# Visual Guide to NoSQL Systems





# Basically Available

Indicates that the system does guarantee availability

# Soft State

Indicates that the state of the system may change due to the eventual consistency

# Eventual Consistency

Indicates that the state of the system will become consistent (if it has the opportunity – is not disturbed)

BASE

---



ACID	BASE
Atomic	Basically Available
Consistency	Soft State
Isolation	Eventual Consistency
Durable	

BASE

---



# Eric Brewer's revision of the CAP theorem and BASE

(Why "2 of 3" is misleading)

Finally, all three properties are more continuous than binary. Availability is obviously continuous from 0 to 100 percent, but there are also many levels of consistency, and even partitions have nuances, including disagreement within the system about whether a partition exists.

<http://www.infoq.com/articles/cap-twelve-years-later-how-the-rules-have-changed>

## Why "2 of 3" is misleading...

---



**Group 1:**

Christian  
Steffen BP  
Troels  
Andreas Bøsig

**Group 2:**

Jakob Bak  
David  
Alice  
Frederik

**Group 3:**

Kaloyan  
Simeon  
Nikolay  
Elias

**Group 4:**

Steffen P  
Mikkel VBA  
Nikolaj  
Rostislav

**Group 5:**

Neli  
Michelle  
Mikkel Ottesen  
Andreas WK

**Group 6:**

Martin  
Jannie  
Mikkel HK  
Terkel

**Group 7:**

Martynas  
Natalia  
Adrian  
Thomas

**Group 8:**

Mike  
Mihail  
Malik  
Vilius

# Groups for mandatory

---



Break 5 minutes

---



- Pick two of your favourite series with at least 4 seasons on OMDb (e.g. Game of Thrones and Modern Family)
- Make a database in your preferred mongo tool called OMDb.
- Insert each season of your series into a collection called series.
- That means that you now should have at least 8 documents in your collection series.

Series database

---



- June 6th, 7th or 8th sometime between 8.30 and 16.00
- English or Danish
- 30 minutes preparation after receiving your question
  - Books, papers, computer, internet allowed
  - No communication with others
    - No calling a friend
    - No Facebook
    - No StackOverflow or other foras
- 20 minutes examination
- 5 minutes for deciding the grade and giving you the reason for it

## About the exam

---



- Data Modeling
- SQL – stored procedures and views
- SQL – transactions and triggers
- JSON
- XML
- NoSQL

Subjects for the exam

---





# STORED PROCEDURES AND VIEWS

```
CREATE PROCEDURE [dbo].[selectTagsFromBlogPost]
-- Add the parameters for the stored procedure here
@blogid int = 0
AS
BEGIN
-- SET NOCOUNT ON added to prevent extra result sets from
-- interfering with SELECT statements.
SET NOCOUNT ON;

    -- Insert statements for procedure here
SELECT T.TagName
FROM Tag as T
INNER JOIN TagBlogPost as TBP ON TBP.TagId = T.TagId
INNER JOIN BlogPost as BP ON BP.BlogPostId = TBP.BlogPostId
WHERE BP.BlogPostId = @blogid;
END
```

What does this code snippet do? How would you use this code?

Explain what conclusions you can draw on the data model behind this. Reflect upon what benefits one may get from placing this in a stored procedure – and are there any cons?

## Exam

---